


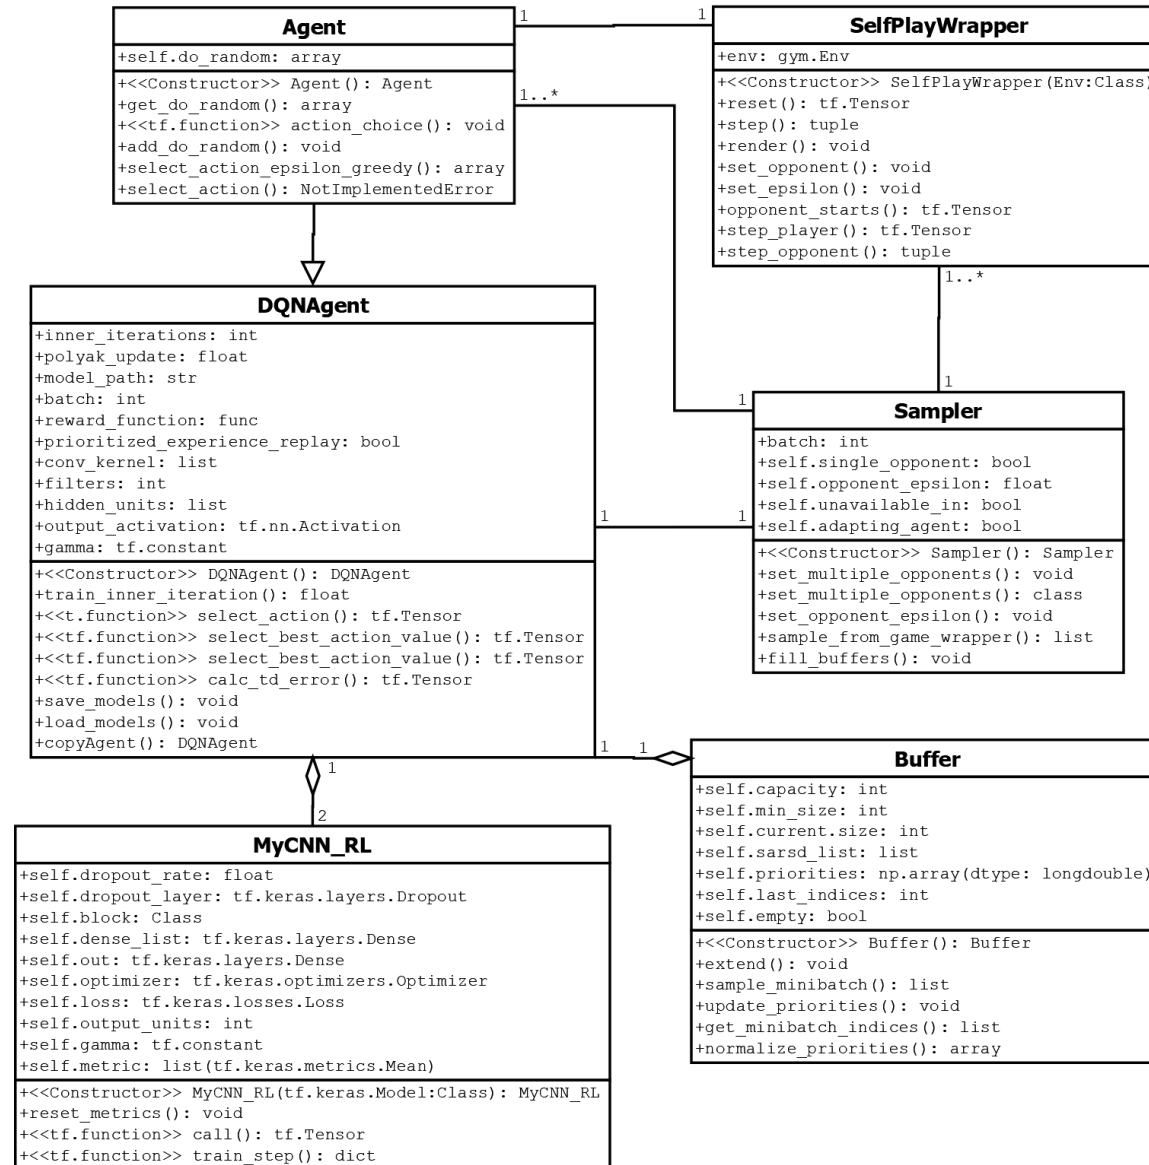


Using Deep Reinforcement learning for creating an adapting agent playing zero-sum games



Project by:
Eosandra Grund
Fabian Kirsch



```
class SelfPlayWrapper(Env):
```

```
    """ A Wrapper for Env similar to keras-gym Connect Four (adapted to a newer python version)
    also works for similar other envs
```

```
    Attributes:
```

```
        env (gym.Env): the gym env to wrap around, has to be a 2 player game
        opponent (Agent): has to be an agent
        epsilon (float): the epsilon value for the epsilon greedy policy
        first_reward: will be used to save the reward, the agent got until the oponent did its step
        last_wrong: will be used to save whether the agent did a wrong move, to make sure the opponent cannot make a move
    """
```

```
    def __init__(self, env_class, opponent = None, epsilon : float = 0):
```

```
        super(Env, self).__init__()
        self.env = env_class()
        self.opponent = opponent
        self.epsilon = epsilon
        self.first_reward = None
        self.last_wrong = None
```

```
        ...
```

```
    def opponent_starts(self):
```

```
        """ let the opponent do the first action, works similar to reset"""
```

```
        s_0 = self.reset()
        # get the opponent's action
```

```
        o_action = self.opponent.select_action_epsilon_greedy(self.epsilon, tf.expand_dims(tf.cast(s_0, dtype = tf.float32), ax
```

```
False)[0]
```

```
        # do the opponent's action
```

```
        s_1,_,_ = self.env.step(o_action, return_wrong = False)
```

```
        return tf.cast(s_1, dtype=tf.float32)
```

```
    def step(self, a, unavailable_in : bool = False, agent = None):
```

```
        # do my step
```

```
        if unavailable_in:
```

```
            s_0, r_0, d_0, w_0 = self.env.step(a, return_wrong = True)
```

```
        else:
```

```
            s_0, r_0, d_0 = self.env.step(a, return_wrong = False)
```

```
            w_0 = False
```

```
        if d_0 or w_0:
```

```
            return tf.cast(s_0, dtype= tf.float32), r_0, d_0
```

```
        # get the opponent's action
```

```
        o_action = self.opponent.select_action_epsilon_greedy(self.epsilon, tf.expand_dims(tf.cast(s_0, dtype = tf.float32)
```

```
False)[0]
```

```
        # do the opponent's action
```

```
        s_1, r_1, d_1 = self.env.step(o_action, return_wrong = False)
```

```
        # calculate the returns
```

```
        if d_1:
```

```
            return tf.cast(s_1, dtype= tf.float32), self.env.loss_reward if r_1 == self.env.win_reward else r_1, d_1
```

```
        return tf.cast(s_1, dtype= tf.float32), r_0, d_1
```

```
class DQNAgent(Agent):
```

```
    """ Implements a basic DQN Algorithm
```

```
    Attributes:
```

```
        model (MyCNN_RL): the model to train
```

```
        target_model (MyCNN_RL): the target_model
```

```
        inner_iterations (int): how many inner iterations to do while training
```

```
        polyak_update (float): how big the polyak update should be
```

```
        model_path (str): where to load and save the models
```

```
        buffer (Buffer): the buffer to save and sample data from
```

```
        batch (int): how big minibatches to sample from the buffer
```

```
        reward_function (func): the function to transform reward gotten from the environment with
```

```
        prioritized_experience_replay (bool): whether to use prioritized_experience_replay
```

```
        conv_kernel (list): a list of the conv_kernels to use to create new model and target_model for the copy
```

```
        filters (int): how many filters to use for the conv layers when making a copy
```

```
        hidden_units (list): list of the hidden_units for dense layers for making a copy
```

```
        output_activation (tf.nn.Activation): function to be the output activation of model and target_model (used for copy)
```

```
        gamma (tf.constant)float: the discount factor for future rewards
    """
```

```
    def __init__(self, env, buffer, batch : int, model_path, polyak_update = 0.9, inner_iterations = 10, reward_function = lambda d,r: r,
        conv_kernel = [3], filters = 128, hidden_units = [64], dropout_rate = 0.5, normalisation : bool = True, prioritized_experience_replay : bool = True,
        gamma = tf.constant = tf.constant(0.99), loss_function = tf.keras.losses.MeanSquaredError(), output_activation = None):
```

```
        super().__init__()
```

```
        # create an initialize model and target_model
```

```
        self.model = MyCNN_RL(conv_kernel = conv_kernel, filters = filters, hidden_units = hidden_units, output_units = env.action_space.n,
```

```
            output_activation = output_activation, loss = loss_function,
```

```
            dropout_rate = dropout_rate, normalisation = normalisation, gamma = gamma)
```

```
        self.target_model = MyCNN_RL(conv_kernel = conv_kernel, filters = filters, hidden_units = hidden_units, output_units = env.action_space.n,
```

```
            output_activation = output_activation, loss = loss_function,
```

```
            dropout_rate = dropout_rate, normalisation = normalisation, gamma = gamma)
```

```
        # build models
```

```
        obs = tf.keras.layers.Input(shape=env.reset().shape)
```

```
        env.close()
```

```
        self.model(obs)
```

```
        self.target_model(obs)
```

```
        self.target_model.set_weights(np.array(self.model.get_weights(), dtype = object))
```

```
        # save other variables as attributes
```

```
        self.inner_iterations = inner_iterations
```

```
        self.polyak_update = polyak_update
```

```
        self.model_path = model_path
```

```
        self.buffer = buffer
```

```
        self.batch = batch
```

```
        self.reward_function = reward_function
```

```
        self.prioritized_experience_replay = prioritized_experience_replay
```

```
        self.conv_kernel = conv_kernel
```

```
        self.filters = filters
```

```
        self.hidden_units = hidden_units
```

```
        self.output_activation = output_activation
```

```
        self.gamma = gamma
```

Project by:

Eosandra Grund

Fabian Kirsch

```
class Buffer:
    """
    Implements a replay buffer for our DQNAgent class. Can use prioritized experience replay.

    Attributes:
        capacity (int): the maximum capacity
        min_size (int): the minimum filling size for starting training
        current_size (int): the current size of the buffer
        sarsd_list (list): contains the data
        priorities (np.array): contains the priorities for the sarsd_list elements with the same index
        last_indices (np.array): contains the indices of the last minibatch, for updating the priorities afterwards
        empty (bool): If the buffer is empty
    """

    def __init__(self, capacity, min_size):
        self.capacity = capacity
        self.min_size = min_size
        self.current_size = 0
        self.sarsd_list = []
        self.priorities = np.array([], dtype=np.longdouble)
        self.last_indices = None
        self.empty = True # is False after adding data the first time

    def extend(self, sarsd):
        """ adds new data to memory buffer """
        """
        Parameters:

        sarsd (list): list of new data samples to add to the buffer. each sample being in the order
            state, action, reward, new_state, done, available_next_action_bool given as a tuple

        """

    def sample_minibatch(self, batch_size):
        """ samples a minibatch from the buffer """
        """
        Parameters:

        batch_size(int) = The sample size pulled from sarsd_list

        """

    Returns:

    sample(list) = A sample pulled from sarsd_list containing [state, action, reward, next_state, done]
    """
```

From training.py

```
def train_self_play_best(agents : list, env_class, batch_size_sampling : int, iterations : int, writers : list, epsilon = 1,
    epsilon_decay : float = 0.9, epsilon_min : float = 0.01, sampling : int = 1, unavailable_in : bool = False, opponent_epsilon = lambda x: (x/2),
    d : int = 20, testing_size : int = 100):
    """
```

```
class Sampler:
    """
    Implements an algorithm to sample from a self-play environment using two different agents

    Attributes:
        envs (list): List of all the environments to sample from
        batch (int): how many environments to sample from at the same time
        agent (Agent): the agent to sample for
        opponent (Agent) or list of Agents: The opponent to use or use several
        single_opponent (bool): whether opponent is a single agent or a list
        opponent_epsilon (float): the epsilon value to use for the epsilon greedy policy of the opponent(s)
        unavailable_in (bool): if True, the sampling agent can decide for unavailable actions and gets a penalty as reward back and the env state does not change
        adapting_agent (bool): whether the sampling agent is an adapting agents that wants information about the rewards other than putting it in the buffer.
            Also there is information about the average rewards of the past and the opponents level (decided by the agent) added to the buffer.
    """

    def sample_from_game_wrapper(self, epsilon : float, save = True):
        """
        samples from env wrappers

        Parameters:
            epsilon (float): the epsilon to use for the epsilon greedy policy of the sampling agent
            save (bool): whether the created samples should be saved in the sampling agents buffer
        """

        sarsd = []
        current_envs = self.envs
        agent_turn = np.random.randint(0,2,(self.batch,))
        observations = np.array([env.opponent_starts() if whether else env.reset() for whether,env in zip(agent_turn,current_envs)])

        for e in range(10000):

            # agent turn
            available_actions = [env.available_actions for env in current_envs]
            available_actions_bool = [env.available_actions_mask for env in current_envs]
            actions = self.agent.select_action_epsilon_greedy(epsilon, observations,available_actions, available_actions_bool, unavailable = self.unavailable_in)

            if self.single_opponent:
                o_0 = np.array([env.step_player(actions[i],self.unavailable_in) for i,env in enumerate(current_envs)]) # only state for opponent input

                # opponent turn
                available_actions = [env.available_actions for env in current_envs]
                available_actions_bool = [env.available_actions_mask for env in current_envs]
                o_actions = self.opponent.select_action_epsilon_greedy(self.opponent_epsilon,o_0,available_actions, available_actions_bool, False)
                results = [env.step_opponent(o_actions[i]) for i,env in enumerate(current_envs)] # new state, reward, done,

            else:
                # here the different opponents in the envs are used
                results = [env.step(actions[i],self.unavailable_in) for i,env in enumerate(current_envs)]


            # if adapting agent give the agent information about the player level (only reward from done envs)
            if self.adapting_agent:
                self.agent.add_game_balance_information([results[i][1] for i in range(len(current_envs)) if results[i][2]])

            # get next available actions, as we also save that in the buffer
            available_actions_bool = [env.available_actions_mask for env in current_envs]


            ...
```

Best Agents

Rewards TikTakToe



Rewards ConnectFour



Project by:

Eosandra Grund

Fabian Kirsch

```
class AdaptingAgent(Agent):
    """
    contains a best agent, whose model outputs the expected future reward for every state-action pair.
    This agent then chooses to use the action with the future reward closest to 0.

    Attributes:
        best_agent (DQNAgent): The agent to use for getting the expected future reward
        self.model: the model of the best_agent
        self.target_model: the target_model of the best_agent
        self.game_balance (numpy.array / list): 1D containing the last reward values
        self.max_balance_length (tf.constant, tf.float32): how many rewards from the past will be saved max
        self.opponent_level (tf.constant): the level of the opponent given by the model (only used in AdaptingDQNAg

    def action_choice(self, probs, available_actions_bool = None):
        """
        returns best action, in this case, that makes the future reward closed to 0

        Parameters:
            probs (tf.Tensor): (batch, model output size) the model output to choose an action from
            available_actions_bool (tf.Tensor): a mask showing which actions are available

        # choose action that makes future game_balance closest to zero
        return tf.argmax(tf.math.abs(probs),axis=-1)
```

Equation 1:

$$\text{best}Q_{sa} = \min \sqrt{Q_{sa}^2}$$

```
class AdaptingAgent5(AdaptingAgent3):
    """
    normalizes the expected future reward and then chooses the value closest to - game_balance instead of just 0 """

    def __init__(self, best_agent : DQNAgent, calculation_value : tf.constant = tf.constant(0.3), game_balance_max : int = 500):
        super().__init__(best_agent,tf.constant(0.3), game_balance_max)
        self.calculation_value = calculation_value

    def action_choice(self, probs, available_actions_bool = None):
        """
        returns best action, in this case, that makes the future reward closed to minus the game balance, but scales the values around -1 and 1 first
        this function removes available actions if they are given.

        Parameters:
            probs (tf.Tensor): (batch, model output size) the model output to choose an action from
            available_actions_bool (tf.Tensor): a mask showing which actions are available

        scaled_around_value = tf.subtract(tf.divide(probs,tf.reduce_max(tf.math.abs(probs))), -self.get_game_balance(tensor=True) * self.calculation_value)

        if available_actions_bool != None:
            scaled_around_value = tf.where(available_actions_bool, scaled_around_value, tf.constant(10.))
        adapting_action = tf.argmin(tf.math.abs(scaled_around_value),axis=-1)

        return adapting_action
```

Equation 5:


$$B_r = \frac{\sum \vec{r}_i}{m}, \quad i \in [-m, \dots, -1]$$

Equation 7:


$$\text{best}Q_{sa} = \min \sqrt{\left(\frac{Q_{sa}}{\max \sqrt{Q_{sa}^2}} - cp \times B_r \right)^2}$$

Tic-Tac-Toe


High-Performance DQN



Adapting Agent with Eq1




Adapting Agent with Eq7, cp=2




Connect Four


High-Performance DQN



Adapting Agent with Eq1



Adapting Agent with Eq7, cp=3



Improvements:



Project by:
Eosandra Grund
Fabian Kirsch